When the Law Doesn’t Exist: AI Hallucinations and the Verification Crisis in Family Court

Hannah Moore - February 15, 2026

Executive Summary

Family courts are increasingly exposed to verification failures stemming from the proliferation of generative AI. Hallucinated case citations, synthetic digital exhibits, and unverified drafting assistance are introducing structural risks into proceedings that rely heavily on credibility and digital artifacts. The core problem is not the use of AI itself, but the absence of standardized verification infrastructure. Without embedded systems for citation validation, metadata preservation, and audit traceability, courts will face escalating burdens in authenticating filings. This article argues for a shift from reactive litigation to proactive, structural verification systems to preserve the integrity of the judicial process.

I. The Emergence of Plausible Legal Fiction

In March 2025, the Appellate Court of Maryland encountered a procedural anomaly in *Chukwuemeka Mezu v. Kristen Mezu*. While reviewing a divorce matter involving a Marital Settlement Agreement, the court discovered a brief replete with case citations that simply did not exist. These authorities appeared legitimate at first glance; they were formatted correctly, attributed to real jurisdictions, and written in the authoritative cadence of traditional jurisprudence. Yet, they were entirely fictional.

The source was generative artificial intelligence. A law clerk had used ChatGPT to search for relevant cases and, unaware of the technology's tendency to "hallucinate," included a list of non-existent authorities in the final filing. The court was forced to issue an Order to Show Cause, diverting significant judicial resources from the merits of the case to determine why the record had been contaminated with "procedural poison."

This event illustrates a broader phenomenon: large language models (LLMs) generate statistically plausible legal text, including fabricated precedents that appear structurally authentic. Unlike traditional legal databases, LLMs do not "find" law; they "predict" it. When such material enters the court record, the damage is procedural:

- Judicial Time Consumption: Clerks and judges must manually verify every citation, slowing an already saturated docket.
- Adversarial Drain: Opposing parties must expend resources to refute "phantom" law that never existed.
- Credibility Collapse: Once a party submits hallucinated law, their entire evidentiary record

becomes suspect.
- Institutional Erosion: Confidence in digital drafting and the reliability of the written record is fundamentally undermined.

II. The Verification Bottleneck and the "Accountant Judge"

Family court judges have long complained about being forced to act as "glorified accountants." In cases involving self-represented litigants (approximately 72% of domestic relations cases according to the National Center for State Courts) judges are often handed shoeboxes of unorganized receipts or spreadsheets of Venmo transactions and asked to calculate child support arrears manually.

Now, generative AI adds a new layer of complexity: the "Verification Bottleneck." Litigants are no longer just asking judges to do the math; they are asking them to authenticate reality. A judge who once spent ten minutes reviewing a bank statement must now decide if a text message screenshot has been subtly altered by an LLM or if a voice memo of a spouse making threats is a high-fidelity deepfake.

The system was not budgeted for this forensic role. Most family courts lack the time and technical capacity to determine the provenance of digital artifacts. When a litigant claims, "I didn't send that text; it was generated by an AI," the judge is left with two conflicting stories and zero metadata to break the tie. This mismatch between technological risk and institutional capacity creates a bottleneck that threatens to paralyze the court's ability to issue timely, fact-based rulings.

III. Why Family Court is Structurally Vulnerable

Family law proceedings amplify AI-related risk due to several inherent factors:

1. Informal Evidence Patterns: Unlike commercial litigation, family court relies heavily on informal communications like texts, social media captures, and GPS logs, all of which are easily manipulated.
2. High Emotional Incentives: The high stakes of custody and safety create intense pressure for parties to "perfect" their evidence or "find" law that supports their position.
3. Resource Asymmetry: Pro se litigants often use free, unverified AI tools as "digital lawyers" because they cannot afford legal counsel or citation-checking software.
4. Compressed Timelines: Family courts operate on high-volume dockets where there is no luxury for the extended evidentiary hearings required to authenticate complex digital media.

In *Zhang v. Chen* (2024), a British Columbia court highlighted that AI "hallucinations" occur at an alarmingly high rate, ranging from 69% to 88% depending on the model. In family court, where a single misquoted standard for the "Best Interests of the Child" can alter a child’s life, this margin of error is unacceptable.

IV. The Cost of Unverified Systems

Absent structural safeguards, AI-related failures produce cascading effects across the justice system. We are seeing an increase in evidentiary disputes and expanded motion practice over authenticity, which directly translates to higher litigation costs for families.

Furthermore, a culture of judicial skepticism is emerging. As judges become more aware of AI's capabilities, they may begin to view all digital submissions with suspicion. While healthy skepticism is necessary, it can unfairly disadvantage self-represented litigants who lack the means to "prove" the authenticity of their genuine digital evidence through expert testimony. The result is not technological progress but increased procedural friction.

V. Structural Requirements for Trusted Digital Evidence

To stabilize family court proceedings in the AI era, verification must be embedded into the infrastructure rather than left to individual diligence. We propose five core requirements:

1. Citation Integrity Controls: Mandatory human certification of authority verification alongside automated cross-referencing within e-filing portals.
2. Immutable Upload Preservation: Court-approved systems must store original digital artifacts without alteration, utilizing hash-based integrity verification (SHA-256) to ensure the file has not been tampered with since upload.
3. Metadata Retention: Systems must preserve original timestamps, file attributes, and modification history to allow for forensic review if authenticity is challenged.
4. Audit Logging: A transparent record of access, edits, and export actions for all digital exhibits.
5. Traceability: Chronological bundling of evidence with verifiable labels and embedded timestamp indicators.

VI. Deep Dive: Metadata Standards for Forensic Verification

For the "Evidence Passport" to function, it must capture and protect specific metadata categories that serve as the "digital fingerprints" of an exhibit.

A. Image and Video Standards (EXIF/IPTC)

The Exchangeable Image File Format (EXIF) is the primary target for forensic scrutiny. A verifiable Evidence Passport must include:

- Temporal Markers: The DateTimeOriginal and SubSecTimeOriginal tags, which record precisely when the image was captured, down to the millisecond.
- Geospatial Provenance: GPSLatitude and GPSLongitude coordinates, cross-referenced with GPSTimeStamp to confirm physical presence.
- Hardware Identifiers: The Make, Model, and SerialNumber of the capturing device. Discrepancies between these and the purported source (e.g., an iPhone 14 photo submitted as a Samsung screenshot) are instant red flags.
- Software Footprints: The Software tag, which reveals if an image was processed by

Adobe Photoshop or mobile AI filters (like "Magic Eraser") before submission.

B. Audio and Multi-Media (BWF/EBUCore)

In custody disputes involving voice recordings, the Evidence Passport relies on Broadcast Wave Format (BWF) or EBUCore standards:

- Originator Reference: Unique identifiers for the recording app or hardware.
- Time Reference: The TimeReference field, which anchors audio to a master clock, essential for detecting "splicing" or AI voice cloning artifacts.
- Coding History: A log of any format conversions, which reveals if a file was compressed or modified to hide AI-generated spectrographic anomalies.

C. Document and Text Provenance (XMP)

For chat logs and PDF filings, Extensible Metadata Platform (XMP) standards provide the necessary audit trail:

- Modification History: The xmpMM:History list, which tracks every save and edit session.
- Derivative Tracking: The xmpMM:DerivedFrom tag, which links a document back to its original source, preventing the submission of "re-typed" or AI-polished summaries as original transcripts.

VII. The Proposed Solution: The "Evidence Passport" System

The Evidence Passport model integrates these standards into a structured submission workflow. Under this model, digital exhibits would be submitted through systems that:

- Automatically Extract and Lock: Capture the original metadata (EXIF/XMP) the moment a file is uploaded.
- Cryptographic Hashing: Generate a unique SHA-256 hash value. This acts as a "seal"—if even one pixel or bit is changed later, the hash will fail, alerting the court to tampering.
- Generate Verifiable Logs: Produce a summary report for the judge detailing the provenance and history of the file.

Courts could then assess evidentiary weight based on provenance transparency. Evidence with a high-fidelity "passport" would be presumed authentic, while evidence without it would face higher scrutiny. This model shifts the burden of verification away from the judge's limited time and onto the technology itself, restoring procedural trust.

VIII. Conclusion: Verification as Foundation

The core issue facing family courts is not whether AI will be used, it already is. The issue is whether courts will operate with unverifiable drafting and evidence systems, or adopt infrastructure that restores procedural trust. Verification architecture cannot be considered an optional upgrade; it is foundational to maintaining the credibility of the law in the AI era. The law must not only be cited; the system must be able to prove it exists.

Citations and References

Chukwuemeka Mezu v. Kristen Mezu, No. 361, September Term, 2025 (Appellate Court of Maryland, Oct. 2025).

Mata v. Avianca, Inc., 678 F. Supp. 3d 443 (S.D.N.Y. 2023).

Zhang v. Chen, 2024 BCSC 285 (Supreme Court of British Columbia, 2024).

National Center for State Courts (NCSC), "Ensuring the Right to be Heard: Guidance for Trial Judges in Cases Involving Self-Represented Litigants" (2025).

Dahl, Matthew, et al. "Large Legal Fictions: Profiling Legal Hallucinations in Large Language Models." arXiv:2401.01301 (2024).

Federal Court of Canada, "Notice to the Profession: Use of Artificial Intelligence in Court Proceedings" (May 2024).

ISO/IEC 27037:2012, "Guidelines for identification, collection, acquisition and preservation of digital evidence."